
RESFIELDS: RESIDUAL NEURAL FIELDS FOR SPATIOTEMPORAL SIGNALS

Marko Mihajlovic¹, Sergey Prokudin^{1,3}, Marc Pollefeys^{1,2}, Siyu Tang¹
ETH Zurich¹; Microsoft²; ROCS, University Hospital Balgrist, University of Zürich ³

markomih.github.io/ResFields

ABSTRACT

Neural fields, a category of neural networks trained to represent high-frequency signals, have gained significant attention in recent years due to their impressive performance in modeling complex 3D data, such as large neural signed distance (SDFs) or radiance fields (NeRFs) via a single multi-layer perceptron (MLP). However, despite the power and simplicity of representing signals with an MLP, these methods still face challenges when modeling large and complex temporal signals due to the limited capacity of MLPs. In this paper, we propose an effective approach to address this limitation by incorporating temporal residual layers into neural fields, dubbed ResFields, a novel class of networks specifically designed to effectively represent complex temporal signals. We conduct a comprehensive analysis of the properties of ResFields and propose a matrix factorization technique to reduce the number of trainable parameters and enhance generalization capabilities. Importantly, our formulation seamlessly integrates with existing techniques and consistently improves results across various challenging tasks: 2D video approximation, dynamic shape modeling via temporal SDFs, and dynamic NeRF reconstruction. Lastly, we demonstrate the practical utility of ResFields by showcasing its effectiveness in capturing dynamic 3D scenes from sparse RGBD cameras of a lightweight capture system.

1 INTRODUCTION

Multi-layer perceptron (MLP) is the most common neural network architecture for representing neural continuous spatiotemporal fields, also known as neural fields (Xie et al., 2022), due to its ability to encode continuous signals over arbitrary dimensions (Kim & Adali, 2003), inherent implicit regularization (Goodfellow et al., 2016; Neyshabur et al., 2014) and spectral bias (Rahaman et al., 2019) that enable good interpolation capabilities. Due to these remarkable properties, MLPs have seen widespread success in many applications such as image synthesis, animation, texture generation, and novel view synthesis (Tewari et al., 2022; Xie et al., 2022).

However, the spectral bias of MLPs (Rahaman et al., 2019), which refers to the tendency of neural networks to learn functions with low frequencies, presents a challenge when it comes to accurately representing complex real-world signals and capturing fine-grained details. Previous efforts have aimed to address the spectral bias by utilizing techniques like positional encoding (Vaswani et al., 2017; Mildenhall et al., 2020; Zhong et al., 2019; Müller et al., 2022) or special activation functions (Sitzmann et al., 2020b; Fathony et al., 2020). However, even with these methods, representing fine-grained details remains a challenge, particularly when dealing with large spatiotemporal signals such as long videos or dynamic 3D scenes.

A straightforward way of increasing the capacity of MLPs is to increase the network complexity in terms of the total number of neurons. However, such an approach would make the inference and optimization slower and more GPU memory expensive, as the time and memory complexity scales with respect to the total number of parameters. Another possibility is to meta-learn MLP weights (Sitzmann et al., 2020a) and maintain specialized independent parameters, but this imposes slow training that does not scale to photo-realistic reconstructions (Tancik et al., 2021). By far the most popular approach for increasing modeling capacity is to partition the spatiotemporal field and fit

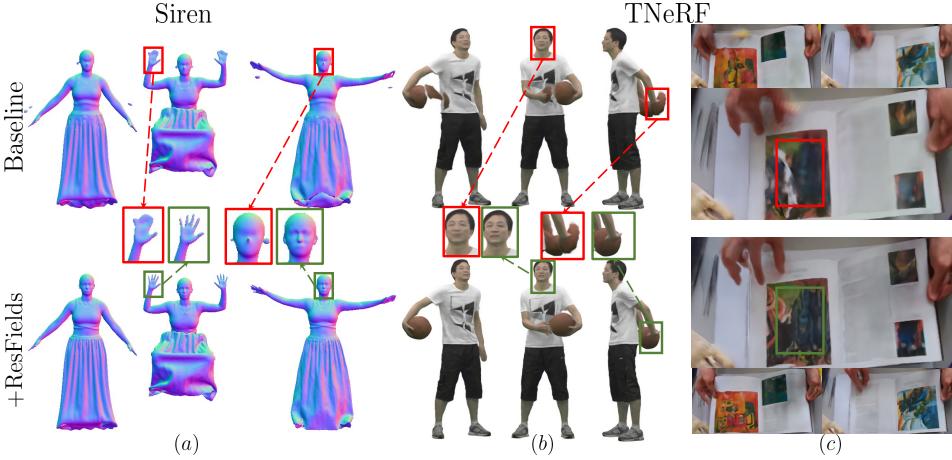


Figure 1: **ResField** extends an MLP architecture to effectively represent complex temporal signals by replacing the conventional linear layers with Residual Field Layers. As such, ResField is versatile and straightforwardly compatible with most existing temporal neural fields. Here we demonstrate its applicability on three challenging tasks by extending Siren (Sitzmann et al., 2020b) and TNeRF (Li et al., 2022): (a) learning temporal signed distance fields and (b) neural radiance fields from four RGB views and (c) from three RGBD views captured by our lightweight rig with three synchronized cameras.

separate/local neural fields (Reiser et al., 2021; Müller et al., 2022; Chen et al., 2022). However, these approaches hinder global reasoning and generalization due to local gradient updates of grid structures (Peng et al., 2023), which is of utmost importance for radiance field reconstruction from sparse views.

ResFields. The challenge that we aim to address in this paper is increasing the model capacity in a way that is agnostic to the design choices of MLP neural fields such as architecture, input encoding, and activation functions. At the same time, we aim to maintain the implicit regularization property of neural networks and complement existing techniques developed for reducing the spectral bias (Mildenhall et al., 2020; Sitzmann et al., 2020b). Our simple key idea is to substitute one or several MLP layers with time-dependent layers (see Fig. 2) whose weights are modeled as trainable residual parameters $\mathcal{W}_i(t)$ added to the existing layer weights \mathbf{W}_i . We dub neural fields implemented in this way ResFields.

Increasing the model capacity in this way offers three key advantages. First, the underlying MLP does not increase in width and hence, maintains the inference and training speed. This property is crucial for most practical downstream applications of neural fields, including NeRF (Mildenhall et al., 2020) which aim to solve inverse volume rendering (Drebin et al., 1988) by querying neural fields billions of times. Second, this modeling retains the implicit regularization and generalization properties of MLPs, unlike other strategies focused on spatial partitioning (Reiser et al., 2021; Müller et al., 2022; Peng et al., 2023). Finally, ResFields are versatile, easily extendable, and compatible with most MLP-based methods for spatiotemporal signals.

However, the straightforward implementation of ResFields could lead to reduced interpolation properties due to a large number of unconstrained trainable parameters. To this end, inspired by well-explored low-rank factorized layers (Denil et al., 2013; Ioannou et al., 2015; Khodak et al., 2021), we propose to implement the residual parameters as a global low-rank spanning set and a set of time-dependent coefficients. As we show in the following sections, this modeling enhances the generalization properties and further reduces the memory footprint caused by maintaining additional network parameters. In summary, our key contributions are:

- We propose an architecture-agnostic building block for modeling spatiotemporal fields that we dub ResFields.

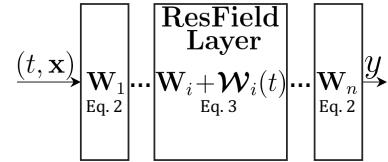


Figure 2: **ResFields**.

-
- We systematically demonstrate that our method benefits a number of existing methods (Sitzmann et al., 2020b; Pumarola et al., 2021; Li et al., 2022; Park et al., 2021b;a; Cai et al., 2022).
 - We validate our residual fields on four challenging tasks and demonstrate state-of-the-art results (Fig. 1): 2D video approximation, temporal 3D shape modeling via signed distance functions, and neural-radiance field reconstruction of dynamic scenes from sparse calibrated RGB and RGBD cameras.

The code, models, and captured data are released at github.com/markomih/ResFields.

2 RELATED WORK

Neural field is a field – a physical quantity that has a value for every point in time and space – that is parameterized fully or in part by a neural network (Xie et al., 2022), typically an MLP as the universal approximator (Kim & Adali, 2003). However, straightforward fitting of signals to regular MLPs yields poor reconstruction quality due to the spectral bias of learning low frequencies (Rahaman et al., 2019). Even though this issue has been alleviated through special input encodings (Mildenhall et al., 2020; Barron et al., 2021; 2022) or activation functions (Sitzmann et al., 2020b; Tancik et al., 2020; Fathony et al., 2020; Lindell et al., 2022; Shekarforoush et al., 2022), neural fields still cannot scale to long and complex temporal signals due to the limited capacity. A natural way of increasing the modeling capacity is by increasing the network’s size in terms of the total number of parameters. However, this trivial solution does not scale well with GPU and training time requirements.

Hybrid neural fields leverage explicit grid-based data structures with learnable feature vectors to improve the modeling capacity via spatial (Takikawa et al., 2021; Müller et al., 2022; Chen et al., 2022; Chan et al., 2022) and temporal (Shao et al., 2023; Fridovich-Keil et al., 2023; Cao & Johnson, 2023; Peng et al., 2023) partitioning techniques. However, these approaches sacrifice the desired global reasoning and implicit regularization property (Neyshabur et al., 2014; Goodfellow et al., 2016) that is needed for generalization, especially for solving ill-posed problems like inverse rendering. In contrast, our proposed solution, ResFields, focuses on improving pure neural network-based approaches that still hold state-of-the-art results across several important applications, as we will demonstrate later.

Input-dependent MLP weights is another common strategy for increasing the capacity of MLPs by directly regressing MLP weights, e.g. via a hypernetwork (Mehta et al., 2021; Wang et al., 2021b) or a convolutional (Peng et al., 2023) neural network. However, these approaches introduce an additional, much larger network that imposes a significant computational burden for optimizing neural fields. KiloNeRF (Reiser et al., 2021) proposes to speed up the inference of static neural radiance fields by distilling the learned radiance field into a grid of small independent MLPs. However, since a bigger MLP is still used during the first stage of the training, this model has the same scaling limitations as the original NeRF. Closest in spirit to our approach, the level-of-experts (LoE) model (Hao et al., 2022) introduces an input-dependent hierarchical composition of shared MLP weights, at the expense of reduced representational capacity. Compared to this work, ResFields demonstrate stronger implicit regularization and representational power for modeling spatiotemporal signals.

Temporal fields are typically modeled by feeding the time-space coordinate pairs to neural fields. SIREN (Sitzmann et al., 2020b) was one of the first neural methods to faithfully reconstruct a 2D video signal. However, scaling this approach to 3D is infeasible and does not produce desired results as demonstrated in dynamic extensions of NeRF models (Pumarola et al., 2021; Li et al., 2022). Therefore, most of the existing solutions (Pumarola et al., 2021; Park et al., 2021a) decouple the learning problem into learning a static canonical neural field and a deformation neural network that transforms a query point from the observation to the canonical space where the field is queried. However, these methods tend to fail for more complex signals due to the difficulty of learning complex deformations via a neural network, as observed by Gao et al. (2022). To alleviate the problem, HyperNeRF (Park et al., 2021b) introduced an additional small MLP and per-frame learnable ambient codes to better capture topological variations, increase the modeling capacity, and simplify the learning of complex deformation. The recent NDR (Cai et al., 2022), a follow-up work of HyperNeRF, further improves the deformation field by leveraging invertible neural networks. All of these methods are fully compatible with the introduced ResFields paradigm which consistently improves baseline results.

Residual connections have a long history in machine learning. They first appeared in Rosenblatt (1961) in the context of coupled perceptron networks. Rosenblatt's insight was that the residual connections increase the efficiency of responding to input signals. Since then, residual connections have been extensively studied and found a major practical utility as a solution to training deep neural networks by overcoming the vanishing gradient problem (Hochreiter, 1998; Srivastava et al., 2015; He et al., 2016) and become a de facto standard for modeling neural networks. Unlike these residual connections that are added to the output of MLP layers, our ResField layers model the residuals of the MLP weights, which in turn yields higher representation power of neural fields, making them more suitable for modeling complex real-world spatiotemporal signals. To the best of our knowledge, directly optimizing residual or multiplicative correctives of model parameters has been explored in the context of fine-tuning large language models (Karimi Mahabadi et al., 2021; Hu et al., 2021; Dettmers et al., 2023) or predicting model weights (Wang et al., 2021b), and has not been explored for directly training spatiotemporal neural fields.

3 RESFIELDS: RESIDUAL NEURAL FIELDS FOR SPATIOTEMPORAL SIGNALS

Formulation. Temporal neural fields encode continuous signals $f : \mathbb{R}^d \times \mathbb{R} \mapsto \mathbb{R}^c$ via a neural network Φ_θ , where the input is a time-space coordinate pair ($t \in \mathbb{R}$, $\mathbf{x} \in \mathbb{R}^d$) and the output is a field quantity $y \in \mathbb{R}^c$. More formally, the temporal neural field is defined as:

$$\Phi_\theta(t, \mathbf{x}) = \sigma_n(\mathbf{W}_n(\phi_{n-1} \circ \phi_{n-2} \circ \dots \circ \phi_1)(t, \mathbf{x}) + \mathbf{b}_n), \quad (1)$$

$$\phi_i(t, \mathbf{x}_i) = \sigma_i(\mathbf{W}_i \mathbf{x}_i + \mathbf{b}_i), \quad (2)$$

where $\phi_i : \mathbb{R}_i^N \mapsto \mathbb{R}_i^M$ is the i th layer of the MLP, which consists of the linear transformation by the weight matrix $\mathbf{W}_i \in \mathbb{R}^{N_i \times M_i}$ and the bias $\mathbf{b}_i \in \mathbb{R}^{N_i}$ applied to the input $\mathbf{x}_i \in \mathbb{R}^{M_i}$, followed by a non-linear activation function σ_i . The network parameters θ are optimized by minimizing a loss term \mathcal{L} directly w.r.t a ground truth signal or indirectly by relating a field quantity to the sensory input, e.g. via volume rendering equation for radiance field reconstruction.

Limitations of MLPs. To model complex and long signals, it is crucial for the underlying MLP to have a sufficient modeling capacity, which scales with the total number of parameters. However, as the MLP size increases, the training time of neural fields becomes slower while increasing the GPU memory requirements, ultimately leading to the bottleneck being the MLP's size. This is especially highlighted for dynamic radiance field reconstruction which requires solving an inverse rendering problem through billions of MLP queries. In the following, we introduce ResFields, an approach for alleviating the capacity bottleneck for spatiotemporal signals.

ResFields model. We introduce residual field layers (Fig. 2) to effectively capture large and complex spatiotemporal signals. ResFields, an MLP that uses at least one residual field layer, alleviates the aforementioned capacity bottleneck without increasing the size of MLPs in terms of the number of layers and neurons. In particular, we replace a linear layer of an MLP ϕ_i with our temporal residual layer defined as:

$$\phi_i(t, \mathbf{x}_i) = \sigma_i((\mathbf{W}_i + \mathcal{W}_i(t))\mathbf{x}_i + \mathbf{b}_i), \quad (3)$$

where $\mathcal{W}_i(t) : \mathbb{R} \mapsto \mathbb{R}^{N_i \times M_i}$ is time-dependent and models residuals of the network weights. This simple formulation increases the model capacity via additional trainable parameters without modifying the overall network architecture.

ResFields factorization. However, naively implementing $\mathcal{W}_i(t) \in \mathbb{R}^{N_i \times M_i}$ as a dictionary of trainable weights would yield a vast amount of independent and unconstrained parameters. This would result in a partitioning of spatiotemporal signal, akin to the space partitioning methods (Reiser et al., 2021; Müller et al., 2022; Shao et al., 2023), and hinder a global reasoning and implicit bias of MLPs, essential properties for solving underconstrained problems such as a novel view synthesis from sparse setups. To this end, inspired by well-established low-rank factorized layers (Denil et al., 2013; Ioannou et al.,

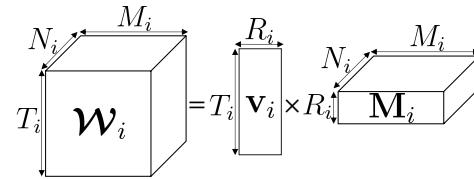


Figure 3: Factorization of \mathcal{W}_i .

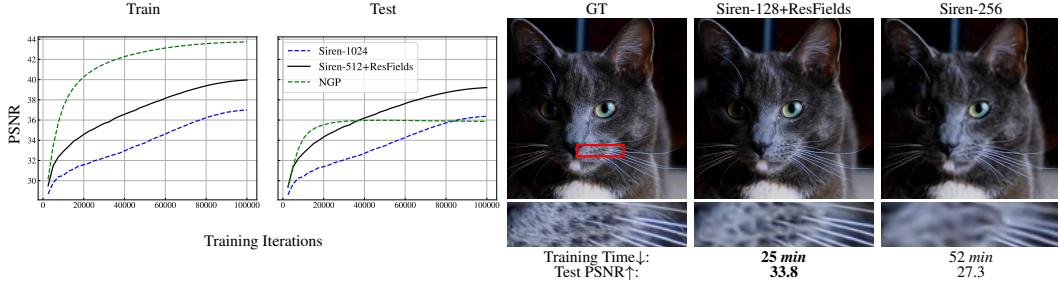


Figure 4: 2D Video Approximation. Comparison of different neural fields on fitting RGB videos. The training and test PSNR curves (left and right respectively) indicate the trade-off between the model’s capacity and generalization properties. Instant **NGP** offers good overfitting capabilities, however, it struggles to generalize to unseen pixels. A Siren MLP with 1024 neurons (**Siren-1024**), shows good generalization properties, however, it lacks representation power (low training and low test PSNR). A smaller Siren with 512 neurons implemented with ResFields (**Siren-512+ResFields**) demonstrates good generalization properties while offering higher model capacity. Besides the higher accuracy, our approach offers ≈ 2.5 times faster convergence and 30% lower GPU memory requirements due to using a smaller network (Tab. 1) Results on the right provide a visual comparison of Siren with 256 neurons and Siren with 128 neurons implemented with ResFields.

2015; Khodak et al., 2021), we directly optimize time-dependent coefficients and R_i -dimensional spanning set of residual network weights that are shared across the entire spatiotemporal signal. In particular, the residual of network weights are defined as

$$\mathcal{W}_i(t) = \sum_{r=1}^{R_i} \mathbf{v}_i(t)[r] \cdot \mathbf{M}_i[r], \quad (4)$$

where the coefficients $\mathbf{v}(t) \in \mathbb{R}^{R_i}$ and the spanning set $\mathbf{M} \in \mathbb{R}^{R_i \times N_i \times M_i}$ are trainable parameters, square brackets denote element selection. To model continuous coefficients over the time dimension, we implement $\mathbf{v} \in \mathbb{R}^{T_i \times R_i}$ as a matrix and linearly interpolate its rows. Such formulation reduces the total number of trainable parameters and further prevents potential undesired overfitting that is common for field partition methods as we will demonstrate later (Sec. 4.4).

Please see the Sup. Mat. for further implementation details.

4 EXPERIMENTS

To highlight the versatility of ResFields, we analyze our method on four challenging tasks: 2D video approximation via neural fields, learning of temporal signed distance functions, and volumetric reconstruction of dynamic scenes from calibrated RGB and RGBD cameras.

4.1 2D VIDEO APPROXIMATION

Learning a mapping of pixel coordinates to the corresponding RGB colors is a popular benchmark for evaluating the implicit model capacity to fit complex signals (Müller et al., 2022; Sitzmann et al., 2020b). For comparison, we use two videos (*bikes* and *cat* from Sitzmann et al. (2020b)) that consist respectively of 250 and 300 frames (with resolutions 512×512 and 272×640) and fit neural representations by minimizing the mean squared error w.r.t ground truth RGB values.

Unlike the proposed setup in Sitzmann et al. (2020b) where the focus is pure overfitting to the image values, our goal is to also evaluate the interpolation aspects of the models. For this, we leave out 10% of randomly sampled pixels for validation and fit the video signal on the remaining ones. We compare our approach against Instant NGP, a popular space grid-based approach to neural field modeling, with the best hyperparameter configuration for the task (see supplementary). We also compare against a five-layer Siren network with 1024 neurons (denoted as Siren-1024), as a pure MLP-based approach. For our model, we choose a five-layer Siren network with 512 neurons, whose hidden layers are implemented as residual field layers with the rank $R_i = 10$ for all hidden layers

(Siren-512+ResFields). We refer to the supplementary for more details and ablation studies on the number of factors, ranks, and layers for the experiment.

Insights. We report both the training and the test PSNR values averaged over the two videos in Fig. 4 and Tab. 1. Here, Instant-NGP offers extremely fast and good overfitting abilities. However, it demonstrates limited generalization to unseen pixels. Siren-1024 has good generalization properties, but clearly underfits the signal and suffers from blur artifacts. Unlike Siren-1024, Siren-512 with our ResFields offers significantly higher reconstruction and generalization quality (≈ 3 PSNR) while requiring 30% less GPU memory and being ≈ 2.5 times faster to train.

This simple experiment serves as a proof of concept and highlights our ability to fit complex temporal signals with smaller MLP architectures, which has a significant impact on the practical downstream applications as we discuss in the following sections.

4.2 TEMPORAL SIGNED DISTANCE FUNCTIONS (SDF)

Signed-distance functions model the orthogonal distance of a given spatial coordinate \mathbf{x} to the surface of a shape, where the sign indicates whether the point is inside the shape. We model a temporal sequence of signed distance functions via a neural field network that maps a time-space coordinate pair ($t \in \mathbb{R}, \mathbf{x} \in \mathbb{R}^3$) to a signed distance value ($y \in \mathbb{R}$).

We sample five sequences of different levels of difficulty (four from Deforming Things (Li et al., 2021) and one from ReSynth (Ma et al., 2021)) and convert the ground-truth meshes to SDF values. We supervise all methods by the MAPE loss following Müller et al. (2022). To benchmark the methods, we extract a sequence of meshes from the learned neural fields via marching cubes (Lorensen & Cline, 1987) and report L1 Chamfer distance (CD \downarrow) and normal consistency (NC \downarrow) w.r.t the ground-truth meshes (scaled by 10^3 and 10^2 respectively). As a main baseline, we use the current state-of-the-art Siren network (five layers) and compare it against Siren implemented with our ResField layers, where residual field layers are applied to three middle layers. We empirically observe that using ResField on the first and last layers has a marginal impact on the performance since weight matrices are small and do not impose a bottleneck for modeling capacity.

Insights. Quantitative and qualitative results (Tab. 2, Fig. 1) demonstrate that ResFields consistently improve the reconstruction quality, with the higher rank increasingly improving results. Importantly, we observe that Siren with 128 neurons and ResFields (rank 40), performs better compared to the vanilla Siren with 256 neurons, making our method over two times faster while requiring less GPU memory due to using a much smaller MLP. Alleviating this bottleneck is of utmost importance for the reconstruction tasks that require solving inverse rendering by querying the neural field billions of times as we demonstrate in the next experiment.

4.3 TEMPORAL NEURAL RADIANCE FIELDS (NERF)

Temporal or Dynamic NeRF represents geometry and texture as a neural field that models a function of color and density. The model is trained by minimizing the pixel-wise error metric between the images captured from known camera poses and ones rendered via the differentiable ray marcher (Mildenhall et al., 2020). To better model geometry, we adopt the MLP architecture and signed distance field formulation from VolSDF (Yariv et al., 2021) that defines density function as Laplace’s cumulative distribution function applied to SDF. We refer the reader to the supplementary for the results with the NeRF backbone and further implementation details.

Following (Wang et al., 2021a), all models are supervised by minimizing the difference between the rendered and the ground truth colors and further adopting the Eikonal (Gropp et al., 2020) and the mask loss terms for well-behaved surface reconstruction under the sparse capture setup:

$$\mathcal{L} = \mathcal{L}_{\text{color}} + \lambda_1 \mathcal{L}_{\text{igr}} + \lambda_2 \mathcal{L}_{\text{mask}}. \quad (5)$$

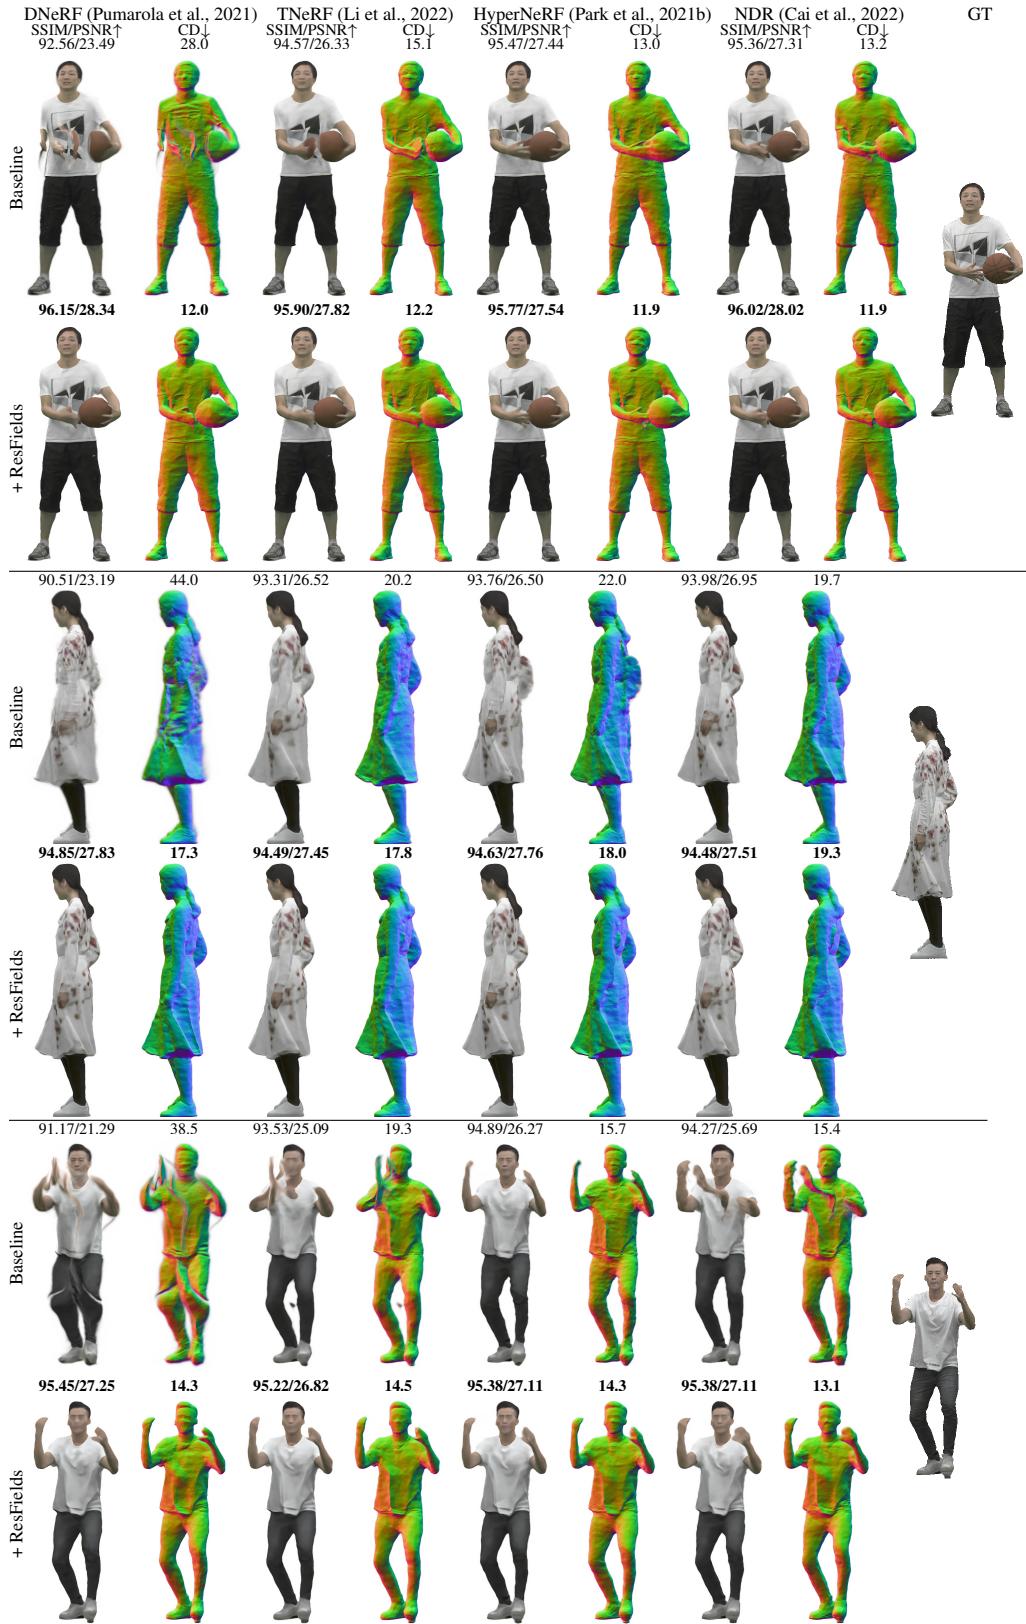


Figure 5: **Temporal radiance field** reconstruction on the Owlii dataset (Xu et al., 2017). Baseline methods implemented with our residual field layers achieve better appearance and geometry reconstruction, with fewer artifacts. This improvement is quantified in Tab. 3; the reported numbers above the images are averaged across all test views of a sequence.

Table 3: **Temporal radiance field** reconstruction on Owlii (Xu et al., 2017). Previous state-of-the-art methods consistently benefit from our residual field layers without imposing a high computational overhead; colors denote the overall 1st, 2nd, and 3rd best-performing model.

	FPS↑	Mean			Basketball			Model			Dancer			Exercise		
		CD↓	SSIM↑	PSNR↑	CD↓	SSIM↑	PSNR↑	CD↓	SSIM↑	PSNR↑	CD↓	SSIM↑	PSNR↑	CD↓	SSIM↑	PSNR↑
Tensor4D (Shao et al., 2023)	0.035	32.8	91.05	22.59	30.5	91.22	22.51	40.3	89.30	22.46	26.7	91.53	23.24	33.5	92.16	22.16
DyNeRF (Li et al., 2022)	0.328	31.0	91.95	23.59	28.0	92.56	23.49	44.9	89.84	23.11	30.7	91.54	23.33	20.3	93.88	24.45
+ ResFields ($i = 1$)	0.327	20.8	93.69	25.57	14.7	94.58	26.54	26.1	92.24	25.36	24.6	93.35	25.20	17.7	94.59	25.17
+ ResFields ($i = 1, \dots, 3$)	0.323	19.3	93.81	25.49	20.3	93.49	24.77	22.2	93.07	26.16	17.6	93.69	25.22	17.1	94.99	25.80
+ ResFields ($i = 1, \dots, 7$)	0.316	19.6	94.00	25.54	17.9	94.47	25.63	23.5	93.15	26.11	20.0	93.58	25.28	16.9	94.81	25.13
TNeRF (Li et al., 2022)	0.339	17.2	94.18	26.18	15.1	94.57	26.33	20.2	93.31	26.52	19.3	93.53	25.09	14.1	95.33	26.77
+ ResFields ($i = 1$)	0.339	14.6	94.99	27.15	12.1	95.67	27.98	18.5	94.07	27.23	14.9	94.59	26.20	13.0	95.63	27.19
+ ResFields ($i = 1, \dots, 3$)	0.334	14.2	95.21	27.44	12.2	95.84	27.98	18.3	94.33	27.81	13.3	94.87	26.55	12.9	95.82	27.40
+ ResFields ($i = 1, \dots, 7$)	0.328	14.2	95.45	27.55	12.2	95.90	27.82	17.8	94.49	27.45	14.5	95.22	26.82	12.3	96.21	28.11
DNeRF (Pumarola et al., 2021)	0.215	32.1	92.09	23.36	22.3	93.21	24.74	44.0	90.51	23.19	38.5	91.17	21.29	23.4	93.47	24.21
+ ResFields ($i = 1$)	0.214	14.2	95.16	27.33	12.1	95.88	28.26	18.1	94.15	27.03	14.1	94.66	26.24	12.8	95.95	27.79
+ ResFields ($i = 1, \dots, 3$)	0.213	14.0	95.34	27.60	12.2	95.95	28.20	17.6	94.45	27.84	14.0	94.88	26.40	12.4	96.08	27.97
+ ResFields ($i = 1, \dots, 7$)	0.210	14.0	95.67	27.89	12.0	96.15	28.34	17.3	94.85	27.83	14.3	95.45	27.25	12.3	96.21	28.14
Nerfies (Park et al., 2021a)	0.180	23.2	93.15	24.35	21.1	93.53	24.74	28.2	92.02	24.25	23.8	92.96	23.81	19.7	94.09	24.60
+ ResFields ($i = 1$)	0.180	14.6	95.12	27.26	12.3	95.64	27.86	19.3	93.95	26.91	14.2	95.10	27.00	12.7	95.77	27.29
+ ResFields ($i = 1, \dots, 3$)	0.179	14.0	95.32	27.43	11.9	95.78	27.87	18.6	94.30	27.21	13.0	95.27	27.11	12.5	95.91	27.51
+ ResFields ($i = 1, \dots, 7$)	0.177	13.8	95.57	27.72	11.8	95.79	27.42	17.6	94.68	27.78	13.5	95.67	27.73	12.2	96.16	27.94
HyperNeRF (Park et al., 2021b)	0.145	16.0	94.94	26.84	13.0	95.47	27.44	22.0	93.76	26.50	15.7	94.89	26.27	13.2	95.64	27.15
+ ResFields ($i = 1$)	0.144	14.4	95.18	27.36	12.4	95.73	28.05	18.7	94.18	27.38	13.4	95.15	26.96	13.0	95.65	27.05
+ ResFields ($i = 1, \dots, 3$)	0.144	14.1	95.35	27.45	12.4	95.86	28.11	18.4	94.36	27.32	12.9	95.35	27.24	12.8	95.82	27.14
+ ResFields ($i = 1, \dots, 7$)	0.143	14.2	95.50	27.64	11.9	95.77	27.54	18.0	94.63	27.76	14.3	95.38	27.11	12.4	96.24	28.16
NDR (Cai et al., 2022)	0.129	15.3	94.82	26.78	13.2	95.36	27.31	19.7	93.98	26.95	15.4	94.27	25.69	12.9	95.65	27.18
+ ResFields ($i = 1$)	0.129	14.7	95.14	27.16	12.6	95.74	27.87	18.2	94.17	27.14	15.2	94.83	26.46	12.9	95.84	27.17
+ ResFields ($i = 1, \dots, 3$)	0.129	14.0	95.36	27.55	12.2	96.00	28.31	18.1	94.29	27.21	13.2	95.12	26.87	12.6	96.04	27.81
+ ResFields ($i = 1, \dots, 7$)	0.127	14.2	95.56	27.81	11.9	96.02	28.02	19.3	94.48	27.51	13.1	95.38	27.11	12.4	96.36	28.61

We use four sequences from the Owlii (Xu et al., 2017) dataset to evaluate the methods. Compared to fully synthetic sequences previously utilised for the task (Pumarola et al., 2021), the dynamic Owlii sequences exhibit more rapid and complex high-frequency motions, making it a harder task for MLP-based methods. At the same time, the presence of ground truth 3D scans allows us to evaluate both geometry and appearance reconstruction quality, as compared to the sequences with only RGB data available (Li et al., 2022; Shao et al., 2023). We render 400 RGB training images from four static camera views from 100 frames/time intervals and 100 test images from a rotating camera from 100 frames. We report L1 Chamfer distance (CD↓) (scaled by 10^3) and the standard image-based metrics (PSNR↑, SSIM↑).

We benchmark recent state-of-the-art methods and their variations with ResFields of rank 10 – TNeRF (Li et al., 2022), DyNeRF (Li et al., 2022), DNeRF (Pumarola et al., 2021), Nerfies (Park et al., 2021a), HyperNeRF (Park et al., 2021b), and NDR (Cai et al., 2022) – as well as a recent timespace-partitioning method Tensor4D (Shao et al., 2023) (with a default training configuration).

Insights. We report all quantitative and qualitative results in Tab. 3 and Fig. 5. Results demonstrate that our method consistently improves all baseline methods, achieving new state-of-the-art results for sparse multi-view reconstruction of dynamic scenes. We further observe that more ResField layers gradually improve results until the point of saturation ($i = 1, \dots, 3$) after which the performance almost levels out since the reconstruction has reached the maximum quality. This experiment confirms that increasing the modeling capacity to a more-than-needed level does not cause overfitting. Importantly, the simplest/cheapest baseline method TNeRF implemented with ResFields performs better than every other more expensive baseline method in the original form. We believe that such speedup and lower memory requirements are of great benefit to the research community, as it enable using lower-end hardware for high-fidelity reconstructions. Given this observation, we set up a simple camera rig and capture longer and more complex sequences to better understand the limitations of ResFields.

Lightweight capture from three RGBD views

RGBD views. We capture four five-second sequences (150 frames) via four synchronized Azure Kinetics (three for reconstruction and one for validation).

Table 4: **Lightweight capture from three RGBD views.**

	Mean		Book		Glasses		Hand		Writing	
	LPIPS↓	SSIM↑	LPIPS	SSIM	LPIPS	SSIM	LPIPS	SSIM↑	LPIPS	SSIM
TNeRF	0.234	79.16	0.323	68.85	0.206	80.44	0.239	81.30	0.168	86.08
+ResFields	0.203	80.00	0.284	70.84	0.164	80.65	0.210	82.09	0.155	86.43

and compare TNerf (w. depth supervision), a baseline with a good balance between computational complexity and accuracy, and its enhancement with ResFields applied to all middle layers. Quantitative evaluation in terms of mean SSIM \uparrow and LPIPS \downarrow (Zhang et al., 2018) reported in Tab. 4 demonstrates that ResFields consistently benefits the reconstruction (see visuals in Fig. 1 and the Sup. video). However, we observe that both methods struggle to capture thin surfaces such as the cord of sunglasses.

4.4 ABLATION STUDY

ResField modeling (Tab. 5). Residual connections on the layer weights ($\mathbf{W}_i + \mathcal{W}_i(t)$) are more powerful compared to modeling residuals on the layer output that is commonly used for conditional generation (Karras et al., 2020), directly modulating layer weights ($\mathbf{W}_i \odot \mathcal{W}_i(t)$) (Mehta et al., 2021), or using time-dependent weights ($\mathcal{W}_i(t)$) as in LoE (Hao et al., 2022). Tab. 5 summarizes the results of these variations on the video approximation task from Sec. 4.1.

Factorization techniques (Tab. 6). We compare our factorization (Eq. 4) with alternative techniques: no factorization (Reiser et al., 2021), hierarchical Levels-of-Experts (LoE) (Hao et al., 2022), and the classic CP Carroll & Chang (1970) and Tucker (1966). CP and Tucker with varying ranks demonstrate good generalization and overfitting results. No factorization achieves great training PSNR, but its generalization performance is sub-optimal which has been mitigated by the hierarchical formulation of LoE. The proposed factorization achieves the best generalization properties. The reported numbers in Tab. 6 are measured on the video approximation task for 30% of unseen pixels. See the Sup. Mat. for additional comparisons.

Limitations. Overall ResFields benefits spatiotemporal neural fields when the bottleneck lies in the modeling capacity rather than in solving unconstrained ill-posed problems. Specifically, we do not observe an advantage on challenging ill-posed monocular reconstruction (Gao et al., 2022) when the main bottleneck is the lack of constraints rather than the network’s capacity.

5 DISCUSSION AND CONCLUSION

We present a novel approach to overcome the limitations of spatiotemporal neural fields in effectively modeling long and complex temporal signals. Our key idea is to incorporate temporal residual layers into neural fields, dubbed ResFields, to enhance the capacity and performance in representing high-frequency signals. The advantage and utility of the method lie in its versatility and straightforward integration into existing works for modeling 2D and 3D temporal fields. ResFields increase the capacity of MLPs without expanding the network architecture in terms of the number of layers and neurons, which allows us to use smaller MLPs without sacrificing the reconstruction quality, while achieving faster inference and training time with a lower GPU memory requirement. We believe that the development towards using lower-cost hardware is the key to democratizing research and making technology more accessible. We hope that the findings of this study contribute to the field of neural fields and provide valuable insights for modeling temporal signals. This, in turn, can lead to advancements in various domains, including computer graphics, computer vision, and robotics.

Table 5: **ResField modeling.**

		Mean PSNR \uparrow	
		test	train
Siren-512	$\phi_i(t, \mathbf{x}_i) = \sigma_i(\mathbf{W}_i \mathbf{x}_i + \mathbf{b}_i)$	31.89	32.13
	+output residual weights (Karras et al., 2020)	- - - - -	- - - - -
	$\phi_i(t, \mathbf{x}_i) = \sigma_i(\mathbf{W}_i \mathbf{x}_i + \mathbf{b}_i) + \mathcal{W}_i(t)$	32.84	33.12
	+modulated weights (Mehta et al., 2021)	- - - - -	- - - - -
	$\phi_i(t, \mathbf{x}_i) = \sigma_i((\mathbf{W}_i \odot \mathcal{W}_i(t)) \mathbf{x}_i + \mathbf{b}_i)$	32.65	32.90
	+direct (Hao et al., 2022) $\mathbf{W}_i(t)$	- - - - -	- - - - -
	$\phi_i(t, \mathbf{x}_i) = \sigma_i(\mathcal{W}_i(t) \mathbf{x}_i + \mathbf{b}_i)$	35.17	35.95
	+ResFields	- - - - -	- - - - -
	$\phi_i(t, \mathbf{x}_i) = \sigma_i((\mathbf{W}_i + \mathcal{W}_i(t)) \mathbf{x}_i + \mathbf{b}_i)$	39.21	39.97

Table 6: **Factorization techniques.**

	Factorization	Rank	#params	Mean PSNR \uparrow	
			[M]	test	train
Siren			0.8	31.96	32.29
	None		236	38.52	48.46
	Reiser et al. (2021)	10 [−]	0.8 [−]	33.04 [−]	33.36 [−]
	CP	20	0.9	33.14	33.47
	Carroll & Chang (1970)	40	1.0	33.41	33.75
		80	1.1	33.72	34.08
		10,64,64 ⁺	1.1 ⁺	33.96 ⁺	34.31 ⁺
		40,64,64	1.5	34.67	35.10
	Tucker (1966)	80,64,64	2.0	35.08	35.59
		10,256,256	3.6	36.31	36.90
		40,256,256	9.5	38.31	39.33
		80,256,256	17.4	39.04	40.39
		(2,4,8) [−]	4.5 [−]	36.42 [−]	37.37 [−]
	LoE	(8,16,32)	15.5	39.87	42.27
	Hao et al. (2022)	(16,32,64)	30.2	40.53	44.15
		(32,64,128)	59.5	40.62	46.35
		10 [−]	8.7 [−]	39.59 [−]	40.80 [−]
	Ours	20	16.5	40.87	42.45
	Eq. 3	40	32.3	41.69	43.72
		80	63.8	41.51	44.39

Acknowledgments and Disclosure of Funding. We thank Hongrui Cai and Ruizhi Shao for providing additional details about the baseline methods and Anpei Chen, Shaofei Wang, and Songyou Peng for proofreading the manuscript. This project has been supported by the Innosuisse Flagship project PROFICIENCY No. PFFS-21-19.

REFERENCES

- Jonathan T Barron, Ben Mildenhall, Matthew Tancik, Peter Hedman, Ricardo Martin-Brualla, and Pratul P Srinivasan. Mip-nerf: A multiscale representation for anti-aliasing neural radiance fields. In *ICCV*, 2021.
- Jonathan T Barron, Ben Mildenhall, Dor Verbin, Pratul P Srinivasan, and Peter Hedman. Mip-nerf 360: Unbounded anti-aliased neural radiance fields. In *CVPR*, 2022.
- Hongrui Cai, Wanquan Feng, Xuetao Feng, Yan Wang, and Juyong Zhang. Neural surface reconstruction of dynamic scenes with monocular rgb-d camera. In *NeurIPS*, 2022.
- Ang Cao and Justin Johnson. Hexplane: A fast representation for dynamic scenes. *CVPR*, 2023.
- J Douglas Carroll and Jih-Jie Chang. Analysis of individual differences in multidimensional scaling via an n-way generalization of “eckart-young” decomposition. *Psychometrika*, 1970.
- Eric R Chan, Connor Z Lin, Matthew A Chan, Koki Nagano, Boxiao Pan, Shalini De Mello, Orazio Gallo, Leonidas J Guibas, Jonathan Tremblay, Sameh Khamis, et al. Efficient geometry-aware 3d generative adversarial networks. In *CVPR*, 2022.
- Anpei Chen, Zexiang Xu, Andreas Geiger, Jingyi Yu, and Hao Su. Tensorf: Tensorial radiance fields. In *ECCV*, 2022.
- Misha Denil, Babak Shakibi, Laurent Dinh, Marc’Aurelio Ranzato, and Nando De Freitas. Predicting parameters in deep learning. *NeurIPS*, 2013.
- Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. Qlora: Efficient finetuning of quantized llms. *arXiv preprint arXiv:2305.14314*, 2023.
- Robert A Drebin, Loren Carpenter, and Pat Hanrahan. Volume rendering. *ACM Siggraph Computer Graphics*, 1988.
- Rizal Fathony, Anit Kumar Sahu, Devin Willmott, and J Zico Kolter. Multiplicative filter networks. In *ICLR*, 2020.
- Sara Fridovich-Keil, Giacomo Meanti, Frederik Rahbæk Warburg, Benjamin Recht, and Angjoo Kanazawa. K-planes: Explicit radiance fields in space, time, and appearance. In *CVPR*, 2023.
- Hang Gao, Ruilong Li, Shubham Tulsiani, Bryan Russell, and Angjoo Kanazawa. Dynamic novel-view synthesis: A reality check. In *NeurIPS*, 2022.
- Xavier Glorot and Yoshua Bengio. Understanding the difficulty of training deep feedforward neural networks. In *AISTATS*, 2010.
- Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep learning*. MIT press, 2016.
- Amos Gropp, Lior Yariv, Niv Haim, Matan Atzmon, and Yaron Lipman. Implicit geometric regularization for learning shapes. In *ICML*, 2020.
- Zekun Hao, Arun Mallya, Serge Belongie, and Ming-Yu Liu. Implicit neural representations with levels-of-experts. *NeurIPS*, 2022.
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Identity mappings in deep residual networks. In *ECCV*, 2016.
- Sepp Hochreiter. The vanishing gradient problem during learning recurrent neural nets and problem solutions. *INT J UNCERTAIN FUZZ*, 1998.

-
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models. In *ICLR*, 2021.
- Yani Ioannou, Duncan Robertson, Jamie Shotton, Roberto Cipolla, and Antonio Criminisi. Training cnns with low-rank filters for efficient image classification. *arXiv preprint arXiv:1511.06744*, 2015.
- Rabeeh Karimi Mahabadi, James Henderson, and Sebastian Ruder. Compacter: Efficient low-rank hypercomplex adapter layers. *NeurIPS*, 2021.
- Tero Karras, Samuli Laine, Miika Aittala, Janne Hellsten, Jaakko Lehtinen, and Timo Aila. Analyzing and improving the image quality of stylegan. In *CVPR*, 2020.
- Mikhail Khodak, Neil Tenenholz, Lester Mackey, and Nicolo Fusi. Initialization and regularization of factorized neural layers. In *ICLR*, 2021.
- Taehwan Kim and Tülay Adalı. Approximation by fully complex multilayer perceptrons. *Neural computation*, 2003.
- Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *ICLR*, 2015.
- Tianye Li, Mira Slavcheva, Michael Zollhoefer, Simon Green, Christoph Lassner, Changil Kim, Tanner Schmidt, Steven Lovegrove, Michael Goesele, Richard Newcombe, et al. Neural 3d video synthesis from multi-view video. In *CVPR*, 2022.
- Yang Li, Hikari Takehara, Takafumi Taketomi, Bo Zheng, and Matthias Nießner. 4dcomplete: Non-rigid motion estimation beyond the observable surface. In *ICCV*, 2021.
- David B Lindell, Dave Van Veen, Jeong Joon Park, and Gordon Wetzstein. Bacon: Band-limited coordinate networks for multiscale scene representation. In *CVPR*, 2022.
- Xiaoxiao Long, Cheng Lin, Peng Wang, Taku Komura, and Wenping Wang. Sparseneus: Fast generalizable neural surface reconstruction from sparse views. *ECCV*, 2022.
- William E Lorensen and Harvey E Cline. Marching cubes: A high resolution 3d surface construction algorithm. *ACM siggraph computer graphics*, 1987.
- Ilya Loshchilov and Frank Hutter. Sgdr: Stochastic gradient descent with warm restarts. *arXiv preprint arXiv:1608.03983*, 2016.
- Qianli Ma, Jinlong Yang, Siyu Tang, and Michael J. Black. The power of points for modeling humans in clothing. In *ICCV*, 2021.
- Ishit Mehta, Michaël Gharbi, Connelly Barnes, Eli Shechtman, Ravi Ramamoorthi, and Manmohan Chandraker. Modulated periodic activations for generalizable local functional representations. In *ICCV*, 2021.
- Ben Mildenhall, Pratul P. Srinivasan, Matthew Tancik, Jonathan T. Barron, Ravi Ramamoorthi, and Ren Ng. Nerf: Representing scenes as neural radiance fields for view synthesis. In *ECCV*, 2020.
- Thomas Müller, Alex Evans, Christoph Schied, and Alexander Keller. Instant neural graphics primitives with a multiresolution hash encoding. *ToG*, 2022.
- Behnam Neyshabur, Ryota Tomioka, and Nathan Srebro. In search of the real inductive bias: On the role of implicit regularization in deep learning. *arXiv preprint arXiv:1412.6614*, 2014.
- Keunhong Park, Utkarsh Sinha, Jonathan T Barron, Sofien Bouaziz, Dan B Goldman, Steven M Seitz, and Ricardo Martin-Brualla. Nerfies: Deformable neural radiance fields. In *ICCV*, 2021a.
- Keunhong Park, Utkarsh Sinha, Peter Hedman, Jonathan T. Barron, Sofien Bouaziz, Dan B Goldman, Ricardo Martin-Brualla, and Steven M. Seitz. Hypernerf: A higher-dimensional representation for topologically varying neural radiance fields. *ToG*, 2021b.

-
- Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, et al. Pytorch: An imperative style, high-performance deep learning library. *NeurIPS*, 2019.
- Sida Peng, Yunzhi Yan, Qing Shuai, Hujun Bao, and Xiaowei Zhou. Representing volumetric videos as dynamic mlp maps. In *CVPR*, 2023.
- Albert Pumarola, Enric Corona, Gerard Pons-Moll, and Francesc Moreno-Noguer. D-nerf: Neural radiance fields for dynamic scenes. In *CVPR*, 2021.
- Nasim Rahaman, Aristide Baratin, Devansh Arpit, Felix Draxler, Min Lin, Fred Hamprecht, Yoshua Bengio, and Aaron Courville. On the spectral bias of neural networks. In *ICML*, 2019.
- Christian Reiser, Songyou Peng, Yiyi Liao, and Andreas Geiger. Kilonerf: Speeding up neural radiance fields with thousands of tiny mlps. In *ICCV*, 2021.
- Frank Rosenblatt. Principles of neurodynamics. perceptrons and the theory of brain mechanisms. Technical report, Cornell Aeronautical Lab Inc Buffalo NY, 1961.
- Ruizhi Shao, Zerong Zheng, Hanzhang Tu, Boning Liu, Hongwen Zhang, and Yebin Liu. Tensor4d: Efficient neural 4d decomposition for high-fidelity dynamic reconstruction and rendering. In *CVPR*, 2023.
- Shayan Shekarforoush, David Lindell, David J Fleet, and Marcus A Brubaker. Residual multiplicative filter networks for multiscale reconstruction. *NeurIPS*, 2022.
- Vincent Sitzmann, Eric Chan, Richard Tucker, Noah Snavely, and Gordon Wetzstein. Metasdf: Meta-learning signed distance functions. *NeurIPS*, 2020a.
- Vincent Sitzmann, Julien Martel, Alexander Bergman, David Lindell, and Gordon Wetzstein. Implicit neural representations with periodic activation functions. *NeurIPS*, 2020b.
- Rupesh Kumar Srivastava, Klaus Greff, and Jürgen Schmidhuber. Highway networks. *arXiv preprint arXiv:1505.00387*, 2015.
- Towaki Takikawa, Joey Litalien, Kangxue Yin, Karsten Kreis, Charles Loop, Derek Nowrouzezahrai, Alec Jacobson, Morgan McGuire, and Sanja Fidler. Neural geometric level of detail: Real-time rendering with implicit 3d shapes. In *CVPR*, 2021.
- Matthew Tancik, Pratul Srinivasan, Ben Mildenhall, Sara Fridovich-Keil, Nithin Raghavan, Utkarsh Singhal, Ravi Ramamoorthi, Jonathan Barron, and Ren Ng. Fourier features let networks learn high frequency functions in low dimensional domains. *NeurIPS*, 2020.
- Matthew Tancik, Ben Mildenhall, Terrance Wang, Divi Schmidt, Pratul P Srinivasan, Jonathan T Barron, and Ren Ng. Learned initializations for optimizing coordinate-based neural representations. In *CVPR*, 2021.
- Ayush Tewari, Justus Thies, Ben Mildenhall, Pratul Srinivasan, Edgar Tretschk, Wang Yifan, Christoph Lassner, Vincent Sitzmann, Ricardo Martin-Brualla, Stephen Lombardi, et al. Advances in neural rendering. In *Computer Graphics Forum*, 2022.
- Ledyard R Tucker. Some mathematical notes on three-mode factor analysis. *Psychometrika*, 1966.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *NeurIPS*, 2017.
- Peng Wang, Lingjie Liu, Yuan Liu, Christian Theobalt, Taku Komura, and Wenping Wang. Neus: Learning neural implicit surfaces by volume rendering for multi-view reconstruction. *NeurIPS*, 2021a.
- Shaofei Wang, Marko Mihajlovic, Qianli Ma, Andreas Geiger, and Siyu Tang. Metaavatar: Learning animatable clothed human models from few depth images. In *NeurIPS*, 2021b.

-
- Yiheng Xie, Towaki Takikawa, Shunsuke Saito, Or Litany, Shiqin Yan, Numair Khan, Federico Tombari, James Tompkin, Vincent Sitzmann, and Srinath Sridhar. Neural fields in visual computing and beyond. In *Computer Graphics Forum*. Wiley Online Library, 2022.
- Yi Xu, Yao Lu, and Ziyu Wen. Owlii dynamic human mesh sequence dataset. In *ISO/IEC JTC1/SC29/WG11 m41658, 120th MPEG Meeting*, 2017.
- Lior Yariv, Jiatao Gu, Yoni Kasten, and Yaron Lipman. Volume rendering of neural implicit surfaces. *NeurIPS*, 2021.
- Richard Zhang, Phillip Isola, Alexei A Efros, Eli Shechtman, and Oliver Wang. The unreasonable effectiveness of deep features as a perceptual metric. In *CVPR*, 2018.
- Ellen D Zhong, Tristan Bepler, Joseph H Davis, and Bonnie Berger. Reconstructing continuous distributions of 3d protein structure from cryo-em images. *arXiv preprint arXiv:1909.05215*, 2019.

A APPENDIX

We provide additional implementation details and experiments to complement our study. All the reported runtime in this paper is measured on an NVIDIA RTX 3090 GPU card.

A.1 IMPLEMENTATION DETAILS

Initialization. For experiments that use Siren networks (Sitzmann et al., 2020b) (sections 4.1 and 4.2), we follow their proposed initialization scheme. Models used for SDF-based dynamic radiance field reconstruction (Sec. 4.3) are initialized following the geometric initialization scheme (Gropp et al., 2020). Other neural network weights are initialized following Glorot & Bengio (2010).

Residual weights. Parameters ($\mathbf{v}_i, \mathbf{M}_i$) which model our residual weights are initialized with a normal distribution $\sim \mathcal{N}(0, 10^{-2})$ to ensure a negligible modification of the initial MLP weights. We observe that larger initial values may negatively affect geometric and Siren initialization. For all experiments in the main paper, we set the number of coefficients T_i to the number of frames. Please see the Sup. Mat. for additional ablation studies and training details.

Training details. All models are trained with Adam optimizer (Kingma & Ba, 2015) with default parameters defined by the PyTorch framework (Paszke et al. (2019)). We observe stable convergence with the learning rate of 5×10^{-4} and gradual cosine annealing (Loshchilov & Hutter, 2016) until the minimum learning rate of 5×10^{-5} for the experiments on dynamic neural radiance fields (Sec. 4.3). For other experiments (sections 4.2 and 4.2), we use the learning rate of 5×10^{-5} and cosine annealing until 5×10^{-6} . All methods are trained respectively for 10^5 , 2×10^5 , 4×10^5 , and 6×10^5 iterations on the 2D video approximation task (Sec. 4.1), temporal SDF reconstruction (Sec. 4.2), and dynamic volumetric reconstruction (Sec. 4.3) on Owlii (Xu et al., 2017) and our captured sequences.

A.2 2D VIDEO APPROXIMATION TASK

All methods presented in the paper (Sec. 4.1) on the 2D video approximation task are trained for 100k iterations, each iteration containing 200k random samples from the training set.

For the NGP baseline in Sec. 4.1, we follow the default setup and use a two-layer fully fused network with ReLU activation functions and run a grid search of hyperparameters to find the optimal configuration. Specifically, we vary the table size T and the number of levels L in Tab. A.5 and found that the best results are achieved with $T = 23$ and $L = 8$ which is reported in the main paper (Fig. 4). Furthermore, we provide results (Tab. A.5) of a much larger five-layer Siren network with 1700 neurons per layer to match the number of trainable parameters to ResFields implemented with a five-layer Siren network with 512 neurons, each containing 8.7M parameters. Expectedly, we observe that both methods achieve similar fitting and generalization performance. However, training such a huge MLP with 1700 neurons becomes impractical, making our approach over six times faster to train while requiring over two times less GPU memory.

Number of factors (Tab. A.1). We further ablate the impact of the number of factors T_i of ResFields. In this experiment, we leave out 10% of randomly sampled pixels for validation and vary the number of factors used for parameterizing the coefficients $\mathbf{v} \in \mathbb{R}^{T_i \times R_i}$, in particular, we set T_i as the percentage of the total number of frames. The results averaged over two videos are reported in Tab. A.1. We observe that the best performance is achieved for 95% when there's little overlap between the coefficients. In practice, there is a negligible difference compared

Table A.1: Number of factors.

Factors T_i	Mean PSNR↑	
	test	train
Siren-512	32.02	32.27
100%	39.86	40.73
95%	39.90	40.77
90%	39.79	40.69
80%	39.69	40.62
70%	39.60	40.49
+ResFields	39.53	40.44
60%	39.45	40.37
50%	39.25	40.20
40%	39.10	40.04
30%	38.87	39.82
20%	38.34	39.29

Table A.2: Time interpolation.

Factors T_i	Mean PSNR↑	
	test	train
Siren-512	26.72	32.36
90 %	21.61	40.90
80 %	22.01	40.82
70 %	24.57	40.76
+ResFields	26.06	40.62
60 %	26.12	40.58
50 %	25.54	40.41
40 %	26.51	40.18
30 %	27.32	39.91
20 %	27.34	39.37

Table A.3: Layers vs. rank.

ResField Layers i	Rank R_i	#params [M]	Mean PSNR↑	
			test	train
2	15	4.7	37.53	38.25
1, 2, 3	5		38.01	38.67
2	30	8.7	38.75	39.69
1, 2, 3	10		39.86	40.73
2	45	12.6	39.33	40.43
1, 2, 3	15		40.62	41.66
2	60	16.5	39.67	40.88
1, 2, 3	20		41.20	42.34

Table A.4: **Ablation study of different fractions of unseen pixels** on the video approximation task. ResFields consistently demonstrate good generalization properties regardless of the difficulty level.

Unseen pixels		Mean PSNR↑		Cat PSNR↑		Bicycles PSNR↑	
		test	train	test	train	test	train
Siren-512	10%	32.02	32.27	31.21	31.41	32.84	33.13
+ ResFields		39.86	40.73	38.58	39.15	41.13	42.32
Siren-1024		36.67	37.36	34.95	35.52	38.38	39.19
+ ResFields		43.15	44.75	42.49	43.53	43.82	45.98
Siren-512	20%	31.99	32.27	31.18	31.41	32.79	33.13
+ ResFields		39.74	40.75	38.50	39.15	40.98	42.35
Siren-1024		36.60	37.39	34.90	35.55	38.30	39.23
+ ResFields		42.95	44.82	42.30	43.53	43.59	46.12
Siren-512	30%	31.97	32.3	31.15	31.42	32.80	33.18
+ ResFields		39.59	40.8	38.39	39.17	40.80	42.43
Siren-1024		36.51	37.42	34.83	35.58	38.19	39.27
+ ResFields		42.72	44.96	42.16	43.60	43.28	46.32
Siren-512	40%	31.91	32.29	31.10	31.41	32.71	33.17
+ ResFields		39.39	40.85	38.28	39.20	40.51	42.50
Siren-1024		36.41	37.50	34.74	35.63	38.08	39.38
+ ResFields		42.33	45.14	41.86	43.67	42.79	46.62
Siren-512	50%	31.85	32.31	31.05	31.42	32.66	33.20
+ ResFields		39.09	40.95	38.08	39.24	40.10	42.66
Siren-1024		36.26	37.61	34.62	35.71	37.90	39.51
+ ResFields		41.75	45.41	41.43	43.79	42.08	47.03
Siren-512	70%	31.59	32.40	30.83	31.48	32.35	33.32
+ ResFields		37.70	41.49	37.14	39.42	38.26	43.55
Siren-1024		35.54	38.04	34.08	36.06	36.99	40.02
+ ResFields		38.96	46.61	39.00	44.44	38.91	48.78

to using independent coefficients (100%) which we use as a default configuration as it is slightly computationally faster.

Time interpolation (Tab. A.2). One downside of using per-frame independent coefficients is that it does not support time interpolation. We conduct an experiment to evaluate the interpolating along the time axis. Here we randomly sample 10% of frames and leave them out for validation. As expected, the lower the number of factors T_i leads to a greater overlap among the frames, consequently leading to a better interpolation properties results, while gradually decreasing the training PSNR.

Layers vs. rank (Tab. A.3). Another natural question to ask is whether it is more beneficial to have more ResField layers or a single ResFied layer with higher ranks while maintaining the constant number of trainable parameters. We conduct this experiment on the video approximation task and compare methods with an equal number of parameters. We conclude that multiple ResField layers provide greater modeling capacity.

Ablation for fewer training samples (Tab. A.4). To complete our study and better understand the implicit bias of our method, we further benchmark Siren and ResFields with varying difficulty levels, ranging from 10-70% of unseen pixels. We observe that ResFields consistently demonstrate good generalization across all the levels of difficulty, well above the baseline.

A.3 TEMPORAL SIGNED DISTANCE FUNCTIONS

The architecture of the Siren MLP used for this experiment is identical to the one used for the video approximation task. All methods are trained for 200k iterations, each batch containing 200k samples uniformly sampled across time. For each frame we follow the sampling strategy for static SDFs (Müller et al., 2022) and sample 50% of points on the mesh, 37.5% normally distributed around the surface $\mathcal{N}(0, 10^{-2})$, and remaining 12.5% are randomly sampled in space.

We provide the full breakdown of per-sequence results in Tab. A.6.

A.4 TEMPORAL NEURAL RADIANCE AND DENSITY FIELDS

All methods on the Owlii dataset are trained for 400k iterations, except Tensor4D for which we follow the default training scheme as the method requires a particular training strategy. The main baselines (TNerf, DyNerf, DNeRF, Nerfies, HyperNeRF, and NDR) are implemented with the MLP architecture from VolSDF (Yariv et al., 2021), however, we reduced the original MLP size

Table A.5: Extended ablation study on the video approximation task. Siren with 512 neurons and residual fields has a total of 8.7M parameters, which is equivalent to Siren with 1700 neurons that indeed achieves similar performance. However, optimizing such a huge MLP comes with a great computational cost with over three times longer training while requiring over two times more GPU memory.

	Resources			Mean		Cat Video		Bikes Video		
	t [it/s]	GPU [G]	#params [M]	test PSNR↑	train PSNR↑	test PSNR↑	train PSNR↑	test PSNR↑	train PSNR↑	
Siren-512	11.66	5.1	0.8	31.89	32.13	31.09	31.29	32.68	32.98	
+ResFields R=10	9.78	6.5	8.7	39.21	39.97	37.96	38.44	40.46	41.50	
Siren-1024	3.55	9.73	3.2	36.37	36.99	34.73	35.24	38.00	38.74	
+ResFields R=10	2.94	13.6	34.6	42.59	43.88	41.88	42.70	43.29	45.05	
Siren-1700	1.42	15	8.7	39.15	40.20	37.26	38.14	41.04	42.25	
NGP T=20	L=6	150	1.3	3.6	31.27	33.25	30.23	31.71	32.31	34.78
	L=7	153	1.3	5.7	32.08	35.06	30.89	33.42	33.28	36.70
	L=8	157	1.2	7.8	32.61	36.64	31.34	35.15	33.88	38.13
	L=9	158	1.1	9.9	32.41	37.80	31.06	36.21	33.75	39.39
NGP T=21	L=6	126	1.6	5.1	32.02	34.48	30.51	32.20	33.52	36.76
	L=7	141	1.5	9.3	32.91	36.87	31.45	34.69	34.37	39.04
	L=8	146	1.3	13.5	33.66	39.33	31.97	37.29	35.34	41.37
	L=9	157	1.2	17.7	33.53	41.36	31.89	39.38	35.17	43.34
NGP T=22	L=6	101	2.1	5.1	32.01	34.47	30.51	32.20	33.51	36.73
	L=7	116	1.7	13.5	33.39	37.82	31.82	35.39	34.96	40.25
	L=8	144	1.5	21.9	34.02	41.09	32.25	38.71	35.79	43.48
	L=9	156	1.2	30.3	33.73	43.83	32.17	41.69	35.29	45.96
NGP T=23	L=6	75	2.7	5.1	32.02	34.48	30.52	32.20	33.52	36.76
	L=7	99	2.1	17.4	33.83	39.51	32.14	36.44	35.52	42.57
	L=8	131	1.6	34.2	34.52	43.85	32.91	40.85	36.13	46.85
	L=9	148	1.2	51	33.96	47.56	32.64	44.63	35.28	50.50
NGP T=24	L=6	51	3.8	5.1	32.01	34.46	30.51	32.20	33.51	36.72
	L=7	77	2.7	17.4	33.84	39.51	32.14	36.45	35.54	42.58
	L=8	130	1.6	51.0	34.27	44.68	32.72	41.71	35.81	47.66
	L=9	145	1.2	84.5	34.02	49.51	33.08	46.58	34.95	52.43

Table A.6: Temporal signed distance function. Siren implemented with our Residual Field layers consistently improves the reconstruction quality. Moreover Siren with 128 neurons and ResFields (rank 40) performs better compared to the much bigger vanilla Siren with 256 neurons. Hence leading to over two times faster inference and convergence time while maintaining lower GPU memory requirements; $t[ms]$ denotes the average inference time for one million query points. Note that using higher ranks almost does not affect the overall time complexity as the main bottleneck is the total number of queries and neurons.

	Resources		Mean		Bear		Tiger		Vampire		Vanguard		ReSynth	
	GPU↓	t [ms]↓	CD↓	ND↓	CD↓	ND↓	CD↓	ND↓	CD↓	ND↓	CD↓	ND↓	CD↓	ND↓
Siren-128	2.4G	20.06	15.063	27.23	7.605	4.519	5.159	4.934	29.490	63.686	17.099	42.953	15.960	20.057
+ ResFields (rank=05)	-	-	9.471	18.537	6.813	3.569	4.422	3.639	12.105	39.457	11.420	30.414	12.594	15.606
+ ResFields (rank=10)	-	-	8.785	16.608	6.671	3.350	4.351	3.435	9.545	32.220	11.140	29.961	12.216	14.075
+ ResFields (rank=20)	2.5G	20.25	8.427	15.483	6.659	3.301	4.325	3.328	8.708	29.661	10.465	27.541	11.980	13.584
+ ResFields (rank=40)	-	-	8.158	14.195	6.579	3.201	4.278	3.148	8.563	29.064	9.729	23.564	11.640	11.999
Siren-256	3.6G	47.99	9.040	16.373	6.532	3.115	4.241	3.055	11.800	36.666	10.623	26.407	12.004	12.622
+ ResFields (rank=05)	-	-	7.901	13.000	6.430	3.009	4.177	2.854	8.576	28.846	8.993	20.168	11.331	10.124
+ ResFields (rank=10)	-	-	7.714	12.242	6.408	3.001	4.161	2.814	8.136	27.404	8.757	18.950	11.110	9.040
+ ResFields (rank=20)	3.8G	48.19	7.662	11.840	6.396	2.995	4.141	2.753	8.076	27.013	8.650	18.056	11.050	8.385
+ ResFields (rank=40)	-	-	7.675	11.674	6.381	3.070	4.137	2.723	8.243	27.296	8.525	17.526	11.087	7.755

Table A.7: **Temporal radiance field reconstruction** on the Owlii dataset (Xu et al., 2017) with the NeRF parametrization. Previous state-of-the-art methods consistently benefit from our residual field layers without the computational overhead; results with the VolSDF parametrization are provided in Tab.3. Colors denote the overall 1st, 2nd, and 3rd best-performing model.

	Mean			Basketball			Model			Dancer			Exercise		
	CD↓	SSIM↑	PSNR↑	CD↓	SSIM↑	PSNR↑	CD↓	SSIM↑	PSNR↑	CD↓	SSIM↑	PSNR↑	CD↓	SSIM↑	PSNR↑
TNeRF (Li et al., 2022)	61.6	93.90	26.04	70.9	94.06	25.88	52.9	93.18	27.06	65.4	93.49	25.22	57.1	94.88	25.99
+ ResFields ($i=1$)	53.6	94.54	26.72	53.8	95.08	27.15	53.4	93.55	27.32	53.3	94.40	26.22	54.0	95.14	26.20
+ ResFields ($i=1, \dots, 3$)	46.1	94.81	27.00	44.6	95.23	27.20	49.6	93.92	27.66	48.1	94.81	26.80	42.3	95.28	26.32
+ ResFields ($i=1, \dots, 7$)	47.6	95.03	27.16	49.6	95.42	27.51	46.3	94.29	27.96	43.5	94.96	26.74	50.9	95.43	26.44
DyNeRF (Li et al., 2022)	60.8	93.68	25.78	60.9	94.03	25.89	56.2	92.60	26.16	70.1	93.46	25.26	56.1	94.63	25.79
+ ResFields ($i=1$)	52.6	94.25	26.42	55.9	94.59	26.70	55.5	93.23	27.08	50.5	94.22	26.13	48.4	94.95	25.75
+ ResFields ($i=1, \dots, 3$)	51.8	94.42	26.53	51.3	94.94	26.74	55.3	93.31	27.07	48.8	94.40	26.33	51.6	95.06	25.99
+ ResFields ($i=1, \dots, 7$)	48.9	94.60	26.72	51.5	95.04	26.95	55.7	93.63	27.40	45.7	94.50	26.23	42.8	95.22	26.28
DNeRF	138.3	92.54	24.40	128.8	92.90	24.34	92.1	91.19	25.04	191.5	92.33	23.50	140.7	93.76	24.72
+ ResFields ($i=1$)	56.5	94.36	26.47	48.0	95.05	27.03	63.8	92.89	26.35	61.9	94.29	26.09	52.2	95.19	26.41
+ ResFields ($i=1, \dots, 3$)	49.6	94.59	26.68	49.3	95.23	27.14	64.0	93.07	26.58	42.2	94.87	26.75	42.8	95.18	26.22
+ ResFields ($i=1, \dots, 7$)	52.7	94.81	26.88	47.5	95.32	26.90	61.6	93.44	26.98	49.0	94.99	26.87	52.5	95.47	26.77
Nerfies (Park et al., 2021a)	135.0	93.57	25.35	97.9	93.55	25.21	135.5	93.10	26.20	186.5	93.41	24.73	120.1	94.23	25.26
+ ResFields ($i=1$)	52.0	94.75	26.99	48.0	95.16	27.14	51.7	93.96	27.79	55.3	94.59	26.36	53.1	95.29	26.66
+ ResFields ($i=1, \dots, 3$)	45.3	94.80	26.88	41.7	95.18	26.70	50.4	93.99	27.82	42.7	94.72	26.49	46.4	95.31	26.52
+ ResFields ($i=1, \dots, 7$)	42.2	94.90	26.73	51.8	95.31	26.71	23.6	93.84	26.97	43.6	95.06	26.76	49.8	95.41	26.49
HyperNeRF (Park et al., 2021b)	63.5	94.67	26.51	59.9	94.64	26.01	57.8	94.25	27.55	69.7	94.44	25.75	66.7	95.35	26.74
+ ResFields ($i=1$)	46.9	94.69	26.80	41.0	95.15	27.14	50.4	93.83	27.70	45.6	94.73	26.54	50.7	95.04	25.80
+ ResFields ($i=1, \dots, 3$)	47.1	94.85	26.99	40.7	95.40	27.41	46.5	93.90	27.71	49.2	94.89	26.75	52.0	95.20	26.07
+ ResFields ($i=1, \dots, 7$)	48.0	95.07	27.27	50.7	95.46	27.50	49.5	94.14	27.90	44.8	95.22	27.03	46.9	95.46	26.66
NDR (Cai et al., 2022)	66.2	94.50	26.48	64.1	94.84	26.64	55.8	94.05	27.40	78.1	93.93	25.34	66.8	95.17	26.53
+ ResFields ($i=1$)	49.5	94.71	26.89	50.4	95.02	26.90	51.8	93.85	27.64	47.6	94.67	26.48	48.3	95.32	26.53
+ ResFields ($i=1, \dots, 3$)	47.5	94.89	27.13	46.2	95.36	27.36	51.0	93.86	27.62	46.4	94.91	26.85	46.1	95.42	26.69
+ ResFields ($i=1, \dots, 7$)	49.8	94.97	27.08	44.2	95.31	27.14	50.7	94.20	27.82	49.3	95.03	26.87	55.2	95.37	26.52

from 256 to 128 neurons as it is impractical to train such large MLPs on more expensive multi-view temporal sequences. The SDF MLP is followed by a two-layer color MLP that takes the output feature of the SDF network. Different from the original color MLP that is conditioned on the viewing ray direction, we do not pass this information to the network as it is impossible to capture any view-dependent appearance for this extremely sparse setup. We observe that training without the viewing direction stabilizes training of the baselines, especially those that rely on a deformation network. We follow the original formulation of the flow MLPs used in DNeRF, Nerfies, HyperNeRF, and NDR.

To make the comparison fair among the baselines, for rendering we employ a non-biased uniform sampling along the ray and sample 1100 rays during the training. On each ray, we sample 256 points where the starting and exiting points of the ray are calculated by ray-box intersection. The box for each sequence is estimated from the ground truth scans with a small padding of 5%. All the methods are supervised by minimizing the loss term in Eq. 5, where we set λ_1 and λ_2 to 0.1.

In practice, the SDF-based density formulation performs better under a sparse setup due to well-behaved surfaces. However, for completeness, we repeat this experiment with the original NeRF formulation (Tab. A.7). The results demonstrate that all of the baselines consistently benefit from ResFields, making Nerfies+ResFields the overall best-performing method in terms of geometry and HyperNeRF+ResFields in terms of appearance.

A.5 PRACTICAL RECONSTRUCTION FROM A LIGHTWEIGHT CAPTURE SYSTEM

We follow the VolSDF architecture from the experiment on the Owlii dataset and train all methods for 600k iterations (2200 rays per batch) since the sequences are longer and images are of higher resolution (540×960).

As the scenes are forward-facing and not constrained from the opposite side, we employ the depth (Cai et al., 2022) and the sparseness loss $\mathcal{L}_{\text{sparse}}$ (Long et al., 2022). In sum, the final loss term is:

$$\mathcal{L} = \mathcal{L}_{\text{color}} + \lambda_1 \mathcal{L}_{\text{depth}} + \lambda_2 \mathcal{L}_{\text{igr}} + \lambda_3 \mathcal{L}_{\text{sparse}}, \quad (\text{A.1})$$

where we set λ_1 , λ_2 , and λ_3 to 0.1 and activate $\mathcal{L}_{\text{sparse}}$ after 70k iterations.



Figure A.1: Data capture.

Data capture. Inspired by the lightweight practical capture of surgeries¹ we create an identical camera rig with four cameras (see Fig. A.1), three for training and one for evaluation. Additionally, we crop captured images by projecting an approximated 3D bounding box around the dynamic region as the main focus of our approach is reconstructing dynamic sequences.

¹OR-X Setup